



# BDI Agents and AgentSpeak(L)

By courtesy of Romelia Plesa  
SITE, University of Ottawa and  
Fariba Sadri, Imperial College London  
Material taken from the Internet

# BDI/AgentSpeak(L)

## References

- A. Rao, AgentSpeak(L): BDI Agents speak out in a logical language, Springer LNCS 1038, 1996
- A. Rao, M. Georgeff, An abstract architecture for rational agents, Proceedings of the 3rd International Conference on Principles of Knowledge Representation and Reasoning, KRR92, Boston, 1992
- R. Bordini et al, Programming MAS in AgentSpeak using Jason, Wiley, 2007

# BDI/AgentSpeak(L)

## Motivations:

- BDI agents are “traditionally” specified in a modal logic with modal operators to represent BDI (Beliefs, Desires and Intentions).
- Their implementations (e.g. PRS, dMARS,jason), however, have typically simplified their specifications and used non-logical procedural approaches.
- AgentSpeak is a programming language based on restricted FOL (First-Order-Logic), similar to prolog.
- AgentSpeak attempts to provide operational and proof theoretic semantics for existing implementations ( and thus by a roundabout way for BDI agents)

# BDI/AgentSpeak(L)

## Further Motivations:

- To incorporate some practical reasoning:
  - Means ends reasoning, deciding how to achieve goals
  - Reaction to events, for example when something unexpected happens
  - Choice deliberation, deciding what we want to achieve (our intention) from amongst our desires

# BDI Agents

- Systems that are situated in a changing environment
- Receive continuous perceptual input
- Take actions to affect their environment

From the various options and alternatives available to it at a certain moment in time, the agent needs to select the appropriate actions or procedures to execute.

The ***selection function*** should enable the system to achieve its objectives, given

- the computational resources available to the system
- the characteristics of the environment in which the system is situated.


# BDI Agents

- two types of input data required for the selection function:
- **Beliefs:**
  - represent the characteristics of the environment
  - are updated appropriately after each sensing action.
  - can be viewed as the ***informative*** component of the system.
- **Desires**
  - contain the information about the objectives to be accomplished, the priorities and payoffs associated with the various objectives
  - can be thought as representing the ***motivational*** state of the system.

# BDI Agents

- **Intentions**
  - represent the currently chosen course of action (the output of the most recent call to the selection function)
  - capture the *deliberative* component of the system.

# BDI Agents



**SELECTION  
FUNCTION**



**INTENTION**





# AgentSpeak(L)

- attempt to bridge the gap between theory and practice
- a model that shows a one-to-one correspondence between the model theory, proof theory and the abstract interpreter.
- natural extension of logic programming for the BDI agent architecture
- provides an elegant abstract framework for programming BDI agents.
- based on a restricted first-order language with events and actions.
- the behavior of the agent (i.e., its interaction with the environment) is dictated by the programs written in AgentSpeak(L).

# AgentSpeak(L) and B-D-Is

→ SPLIT IN PROLOG

- An agent's **belief state** is the current state of the agent, which is a model of itself, its environment, and other agents.
- The agent's **desires** are the states that the agent wants to bring about based on its external or internal stimuli.
- When an agent commits to a particular set of plans to achieve some goal, these partially instantiated plans are referred to as an **intention** associated with that goal. Thus, **intentions** are active plans that the agent adopts in an attempt to achieve its goals.

# AgentSpeak(L) - Basic Notions

- The specification of an agent in AgentSpeak(L) consists of:
  - a set of base ***beliefs***
    - Facts and rules in the logic programming sense
  - a set of ***plans***.
    - context-sensitive, event-invoked recipes that allow hierarchical decomposition of goals as well as the execution of actions with the purpose of accomplishing a goal.

# BDI/AgentSpeak(L)

## Internal (Mental) State

- A set of *beliefs*
- A set of current *desires* (or goals)
  - typically of the form !b where b is belief
  - interpreted as desire for state of the world in which b holds.
- ✱ A set of pending *events*
  - typically perceptions of messages interpreted as belief updates: +b, -b or as goals to be achieved: +!b
  - including request messages from other agents usually recorded as new belief events, perhaps as a new belief that the request has been made.
- A set of *intentions*
- A *plan library*. A plan has a triggering condition (an event), a mental state applicability condition, and a collection of sub-goals and actions (similar to ECA rules).



L'EVENTO PUÒ DIRE:  
- ADD A BELIEVES  
- REMOVE " " "  
- SETWILL GOAL

# Agent nigga (L) Beliefs nigga

## Terms

- No modal operators
- Beliefs: a conjunction of ground literals
- `adjacent(room1, room2) & loc(room1) & ¬empty(room1)`
- Events: If  $b$  is an atomic belief then the following are event terms:
  - $!b$  represents an achievement goal, e.g. `!loc(room2)`
  - $?b$  represents a test goal, e.g. `?empty(room1)`
  - $+b$ ,  $-b$  representing events of adding or deleting beliefs (events generated by messages)
  - $+!b$ ,  $-!b$
  - $+?b$ ,  $-?b$
- Agent can have explicit goals, given by events

# AgentSpeak(L) - Basic Notions

- ***belief atom***
  - is a first-order predicate in the usual notation
  - belief atoms or their negations are termed ***belief literals***.

# AgentSpeak(L) - Basic Notions

- **goal**
  - is a state of the system, which the agent wants to achieve.
- two types of goals:
  - **achievement goals:** ground prolog-like atoms
    - predicates prefixed with the operator “!”
    - state that the agent wants to achieve a state of the world where the associated predicate is true.
    - in practice, these initiate the execution of *subplans*.
  - **test goals:** *prolog-like atoms (possibly non-ground)*
    - predicates prefixed with the operator ‘?’
    - returns a unification for the variables occurring in the atom with one of the agent’s beliefs; it fails if no unification is found.

# AgentSpeak(L) - Basic Notions

- ***triggering event***
  - defines which events may initiate the execution of a plan.
  - an *event* can be
    - internal, when a subgoal needs to be achieved
    - external, when generated from belief updates as a result of perceiving the environment.
  - two types of triggering events:
    - related to the *addition* ('+') and *deletion* ('-') of attitudes (beliefs or goals).



# AgentSpeak(L) - Basic Notions

## Plans

- refer to the *basic actions* that an agent is able to perform on its environment.

$$p ::= te : ct \leftarrow h$$

Where:

- te* - triggering event (denoting the purpose for that plan)
- ct* - a conjunction of belief literals representing a context.
  - The context must be a logical consequence of that agent's current beliefs for the plan to be applicable.
- h* - a sequence of basic actions or (sub)goals that the agent has to achieve (or test) when the plan, if applicable, is chosen for execution.

Triggering  
event

Context

```
+concert (A,V) : likes(A) <-  
    !book_tickets(A,V).
```

Achievement  
goal added

```
+!book_tickets(A, V) :  
    ¬busy(phone)  
    <- call(V);  
    ...;  
    !choose  
seats(A,V) .
```

Basic action

# AgentSpeak(L) - Basic Notions

- ***Intentions***

- plans the agent has chosen for execution.

- Intentions are executed one step at a time.

- A step can

- query or change the beliefs

- perform actions on the external world

- suspend the execution until a certain condition is met

- submit new goals.

- The operations performed by a step may generate new events, which, in turn, may start new intentions.

- An intention succeeds when all its steps have been completed. It fails when certain conditions are not met or actions being performed report errors.

# AgentSpeak Plans

- Each agent has its own repertoire of (primitive) actions and plan library.

Plans are ECA(Event-Condition-Action) rules of the form:

$e:b_1, \dots, b_m \leftarrow h_1; \dots; h_k$

$e$  is an event term

the  $b_i$  are belief terms –  $b_1, \dots, b_m$  is called **context**

the  $h_i$  are goals or (primitive) actions

Plans are used to respond to belief update events and new goal events

# AgentSpeak Plan Example

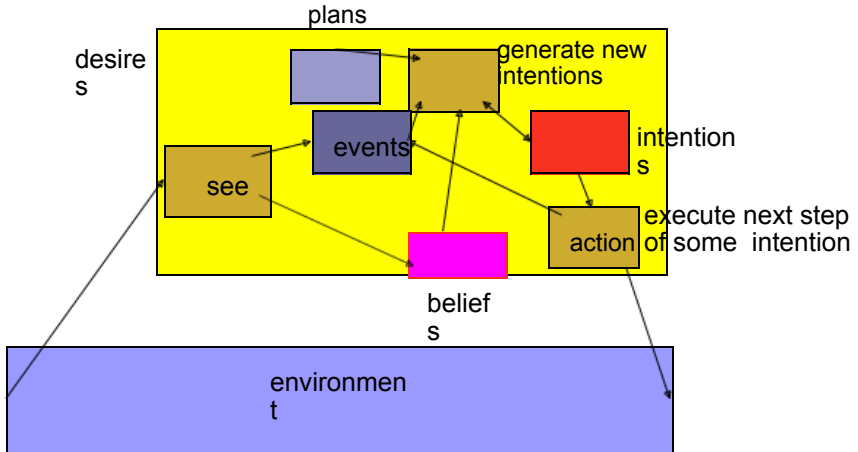
CONTEXT  
+!quench\_thirst:have\_glass <-  
    !have\_soft\_drink; → sotto goal  
    fill\_glass, drink → azioni

+!have\_soft\_drink:soft\_drink\_in\_fridge  
    <- open\_fridge;  
    get\_soft\_drink

# AgentSpeak Plan Example (cunt'd)

```
+!have_soft_drink:not soft_drink_in_fridge  
    <- go_to_shop;  
        buy_soft_drink  
+!have_soft_drink:not soft_drink_in_fridge  
    <- go_to_neighborhood;  
        ask_soft_drink
```

# AgentSpeak Agent Cycle



# AgentSpeak Agent Cycle

1. Notice external/internal changes
2. Update belief and record as events in event stores  
e.g. `+!location(robot, b)`, `+location(waste, a)`
3. **Choose** event (from event store) or desire (from desire store for which there is at least one plan)
4. **Select** plan – this becomes new intention
5. Drop intentions no longer believed viable
6. Resume intention
  - Execute an action, or
  - Post subgoal as a new goal event
7. Repeat cycle



# AgentSpeak(L) Syntax

```
ag ::= bs ps
bs ::= at1. ... atn. (n ≥ 0)
at ::= P(t1, ... tn) (n ≥ 0)
ps ::= p1 ... pn (n ≥ 1)
p ::= te : ct <- h.
te ::= +at | -at | +g | -g
ct ::= true | l1 & ... & ln (n ≥ 1)
h ::= true | f1 ; ... ; fn (n ≥ 1)
l ::= at | not (at)
f ::= A(t1, ... tn) | g | u (n ≥ 0)
g ::= !at | ?at
u ::= +at | -at
```

# AgentSpeak(L)- Informal Semantic

- The interpreter for AgentSpeak(L) manages
  - a set of *events*
  - a set of *intentions*
  - three *selection functions*.

# AgentSpeak(L)- Informal Semantic Events

- **Events**, which may start off the execution of plans that have relevant triggering events, can be:

- **external**, when originating from perception of the agent's environment (i.e., addition and deletion of beliefs based on perception are external events). External events may create new intentions.

- **internal**, when generated from the agent's own execution of a plan (i.e., a subgoal in a plan generates an event of type "addition of achievement goal").

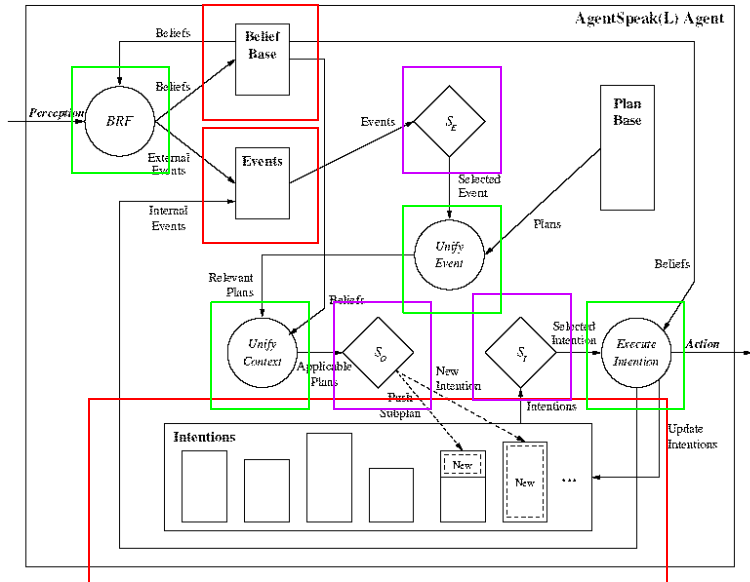
# AgentSpeak(L)- Informal Semantic Intentions

- ***Intentions*** are particular courses of actions to which an agent has committed in order to handle certain events. Each intention is a stack of partially instantiated plans

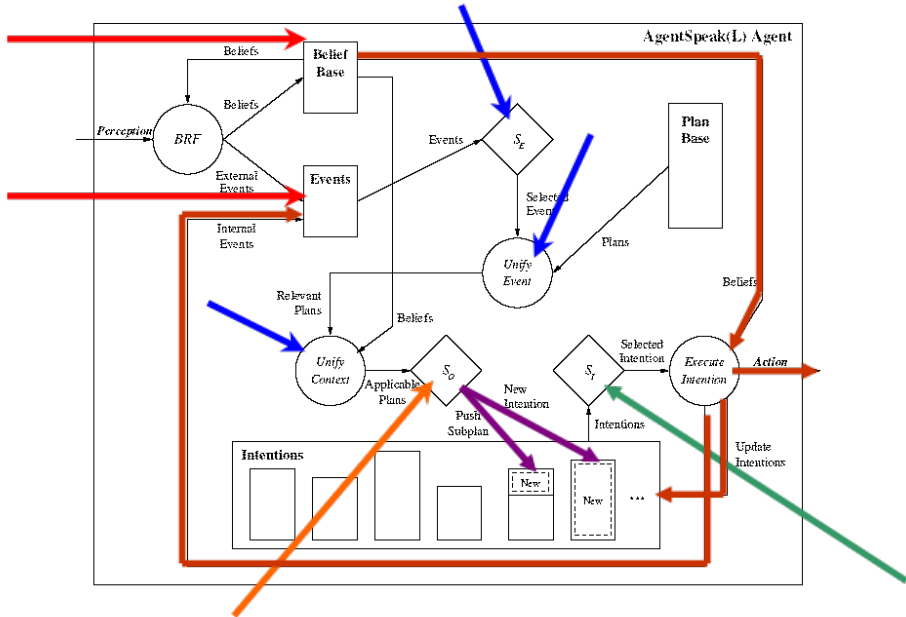
# AgentSpeak(L)- Informal Semantic Selection Functions

- **SE** (the event selection function)
  - selects a single event from the set of events
- **SO**
  - selects an “option” (i.e., an applicable plan) from a set of applicable plans
- **SI**
  - selects one particular intention from the set of intentions.
- *The selection functions are agent-specific, in the sense that they should make selections based on an agent's characteristics.*

# AgentSpeak(L)- Informal Semantic



# AgentSpeak(L)- Informal Semantic



# AgentSpeak(L) Example



ALICE

- During lunch time, forward all calls to Carla.
- When I am busy, incoming calls from colleagues should be forwarded to Denise.



# AgentSpeak(L) Example Beliefs

```
user(alice) .  
user(bob) .  
user(carla) .  
user(denise) .  
~status(alice, idle) .  
status(bob, idle) .  
colleague(bob) .  
lunch_time("11:30") .
```

# AgentSpeak(L) Example Plans

```
user(alice).  
user(bob).  
user(carla).  
user(denise).  
~status(alice, idle).  
status(bob, idle).  
colleague(bob).  
lunch_time("11:30").
```

*"During lunch time, forward all calls to Carla".*

```
+invite(X, alice) : lunch_time(t)    []  
                    !call_forward(alice, X, carla). (p1)
```

*"When I am busy, incoming calls from colleagues should be forwarded to Denise".*

```
+invite(X, alice) :  
    colleague(X)    []  
    !call_forward_busy(alice, X, denise).  
                                (p2)
```

```
+invite(X, Y) : true    []    connect(X, Y).  
                                (p3)
```

# AgentSpeak(L) Example Plans

```
user(alice).  
user(bob).  
user(carla).  
user(denise).  
~status(alice, idle).  
status(bob, idle).  
colleague(bob).  
lunch_time("11:30").  
+invite(X, alice) : lunch_time(t)    [] !call_forward(alice, X, carla).    (p1)  
+invite(X, alice) :      colleague(X) [] call_forward_busy(alice,X,denise).(p2)  
+invite(X, Y): true    [] connect(X,Y).                                (p3)
```

```
+!call_forward(X, From, To) : invite(From, X)  
    [] +invite(From, To), - invite(From,X)  
(p4)
```

```
+!call_forvard_busy(Y, From, To) : invite(From, Y) &  
    not(status(Y, idle))  
    [] +invite(From, To), - invite(From,Y).      (p5)
```

# AgentSpeak(L) Example

```
user(alice).
user(bob).
user(carla).
user(denise).
~status(alice, idle).
status(bob, idle).
colleague(bob).
lunch_time("11:30").

+invite(X, alice) : lunch_time(t)
                    [] !call_forward(alice, X, carla).                (p1)
+invite(X, alice) :      colleague(X)
                    [] call_forward_busy(alice,X,denise).            (p2)
+invite(X, Y): true   [] connect(X,Y).                                (p3)
+!call_forward(X, From, To) : invite(From, X)
    [] +invite(From, To), - invite(From,X)
    (p4)
+!call_forward_busy(Y, From, To) : invite(From, Y)&
    not(status(Y, idle))
    [] +invite(From, To), - invite(From,Y).                            (p5)
```

# Execution - 1

- a new event is sensed from the environment,  
`+invite(bob, alice)` (there is a call for Alice from Bob).
- There are three *relevant* plans for this event (p1, p2 and p3)
  - the event matches the triggering event of those three plans.

Relevant Plans	Unifier
p1: <code>+invite(X, alice) : lunch_time(now)</code> <code>□ !call_forward(alice, X, carla)</code>	
p2: <code>+invite(X, alice) : colleague(X)</code> <code>□ !call_forward_busy(alice, X, denise).</code>	<code>{X=bob}</code>
p3 : <code>+invite(X, Y) : true □ connect(X,Y).</code>	<code>{Y=alice, X=bob}</code>

# Execution - 2

- only the context of plan p2 is satisfied - `colleague(bob)` => p2 is *applicable*.
- a new intention based on this plan is created in the set of intentions, because the event was external, generated from the perception of the environment.
- The plan starts to be executed. It adds a new event, this time an internal event: `!call_forward_busy(alice,bob,denise)` .

Intention ID	Intension Stack	Unifier
1	<code>+invite(X,alice):colleague(X)</code> <code>&lt;- !call_forward_busy(alice,X,denise)</code>	<code>{X=bob}</code>

## Execution - 3

- a plan relevant to this new event is found (p5):

Relevant Plans	Unifier
<pre>p5:  +!call_forward_busy(Y, From, To) :       invite(From, Y) &amp; not(status(Y, idle))       [] +invite(From, To),         - invite(From,Y) .</pre>	<pre>{From=bob, Y=alice, To=denise}</pre>

- p5 has the context condition true, so it becomes an *applicable* plan and it is pushed on top of *intention 1* (it was generated by an internal event)

Intention ID	Intension Stack	Unifier
1	<pre>+!call_forward_busy(Y,From,To) : invite(From,Y) &amp; not status(Y,idle) &lt;- +invite(From,To) ; -invite(From,Y)</pre>	<pre>{From=bob, Y=alice, To=denise}</pre>
	<pre>+invite(X,alice) : colleague(X) &lt;- !call_forward_busy(alice,X,denise)</pre>	<pre>{X=bob}</pre>

# Execution - 4

- A new internal event is created, `+invite(bob, denise)`.
- three relevant plans for this event are found, p1, p2 and p3.
- However, only plan p3 is applicable in this case, since the others don't have the context condition true.
- The plan is pushed on top of the existing intention.

Intention ID	Intension Stack	Unifier
1	<code>+invite(X,Y) : &lt;- connect(X,Y)</code>	<code>{Y=denise, X=bob}</code>
	<code>+!call_forward_busy(Y,From,To) : invite(From,Y) &amp; not status(Y,idle) &lt;- +invite(From,To); -invite(From,Y)</code>	<code>{From=bob, Y=alice, To=denise}</code>
	<code>+invite(X,alice) : colleague(X) &lt;- !call_forward_busy(alice,X,denise)</code>	<code>{X=bob}</code>



# Execution - 5

- on top of the intention is a plan whose body contains an action.
- the action is executed, `connect(bob, denise)` and is removed from the intention.
- When all formulas in the body of a plan have been removed (i.e., have been executed), the whole plan is removed from the intention, and so is the achievement goal that generated it.

Intention ID	Intension Stack	Unifier
1	<code>+!call_forward_busy(Y,From,To)</code> <code>invite(From,Y) &amp; not status(Y,idle)</code> <code>&lt;- -invite(From,Y)</code>	{From=bob, Y=alice, To=denise}
	<code>+invite(X,alice) : colleague(X)</code> <code>&lt;-</code> <code>!call_forward_busy(alice,X,denise)</code>	{X=bob}

- The only thing that remains to be done is `-invite(bob, alice)` (this event is removed from the beliefs base).
- This ends a cycle of execution, and the process starts all over again, checking the state of the environment and reacting to events.

# AgentSpeak Plans

## Another Example: Cleaning Robot

+location(waste, X) :

location(robot,X) & location(bin,Y) & X != Y <-

pick(waste);

!location(robot,Y);

drop(waste).

+location(waste, X) :

location(robot,X) & location(bin,X) <-

pick(waste);

drop(waste).

# AgentSpeak Cleaning Robot

```
+location(waste, X) :  
    location(robot,X) & location(bin,Y) <-  
        pick(waste);  
    !location(robot,Y); drop(waste).
```

Context

Triggering  
Event-  
Addition  
of a fact

Body of the plan

# AgentSpeak Cleaning Robot

```
+location(waste, X) :  
    location(robot,X) & location(bin,Y) <-  
        pick(waste);  
    !location(robot,Y); drop(waste).
```


The intended reading of this is very similar to event-condition-action rules (except that the action part is more sophisticated):

On event of noticing waste at X, if robot is at X and bin at Y, then (robot) pick waste, make its location Y and drop waste.

# AgentSpeak Cleaning Robot

```
+!location(robot, X) :  
    location(robot,X) <- true.
```

```
+!location(robot, X) :  
    location(robot,Y) & not X=Y &  
    adjacent(Y,Z) & not location(robot,Z) <-  
        !move(robot,Y,Z);  
    +!location(robot,X).
```



```
+!move(robot, X,Y) :  
  location(robot,Y) & not X=Y &  
  adjacent(Y,Z) <-  
    -location(robot,Y),  
    moveA(Y,Z),  
    +location(robot,Z),
```

# AgentSpeak Cleaning Robot

```
+!location(robot, X) :  
    location(robot,Y) & not X=Y &  
    adjacent(Y,Z) & not location(robot, Z) <-  
        move(Y,Z); +!location(robot,X).
```

The intended reading of this is similar to  
goal reduction rules:

To achieve a goal `location(robot,X)` ....

# Jason

- a fully-fledged interpreter for AgentSpeak(L)
- many extensions, providing a very expressive programming language for agents.
- allows configuration of a multi-agent system to run on various hosts.
- implemented in Java (thus it is multi-platform)
- available *Open Source* and is distributed under GNU LGPL.
- <http://jason.sourceforge.net/>



# Jason characteristics

- support for developing Environments (Java)
- the possibility to run a multi-agent system distributed over a network
- fully customisable (in Java) selection functions, trust functions, and overall agent architecture (perception, belief-revision, inter-agent communication, and acting)
- a library of essential "internal actions"
- straightforward extensibility by user-defined internal actions, which are programmed in Java.



# Conclusion

- AgentSpeak(L) has many similarities with traditional logic programming, which would favor its becoming a popular language
- it proves quite intuitive for those familiar with logic programming.
- it has a neat notation, thus providing quite elegant specifications of BDI agents.