

Envy-free Scheduling with unrelated Machines

Resources

- As far as the exam is concerned, it is enough studying (and then deeply understanding) the following slides.
- However, if interested, you can find results contained in these slides and much more at this paper:

V. Bilò, A. Fanelli, M. Flammini, G. Monaco, L. Moscardelli: The Price of Envy-Freeness in Machine Scheduling. *Theoretical Computer Science*. Vol. 613, pp. 65—78, 2016.

Introduction

- We consider the Envy-free Scheduling Unrelated Machine MAKESPAN (minimization) problem.
- We still suppose that in any schedule \mathcal{S} all the machines get at least one job, (we can always remove machines getting no job).

k -Envy-free Scheduling Unrelated Machine **MAKESPAN (minimization) problem.**

- **INPUT:** m unrelated machine ($h = 1, \dots, m$), n jobs ($i = 1, \dots, n$), $p_{ih} > 0$, parameter (positive integer) $k \geq 1$.
- **OUTPUT:** k -envy free schedule $S = (S_1, \dots, S_m)$ such that $\sum_{i \in S_j} p_{ij} \leq k \sum_{i \in S_{j'}} p_{ij}$ for any machines j, j' .
- **GOAL:** Minimizing the Machine MAKESPAN $\max_{h=1, \dots, m} \{MC_h(S)\}$
- In other words, a schedule S is k -envy-free if, for any machine j , the completion time of machine j in S is at most k times than the completion time machine j could achieve by getting the jobs of any another machine j' (that gets at least one job in S), for a given factor $k \geq 1$.
- If for some machine j , there exists a machine j' such that $\sum_{i \in S_j} p_{ij} > k \sum_{i \in S_{j'}} p_{ij}$ we will say that machine j *envies* machine j' in the schedule S .

An example

- INPUT: 3 Unrelated machines, 4 jobs.

$p_{ih} =$

$M \backslash J$	J_1	J_2	J_3	J_4
M_1	1	2	3	4
M_2	3	5	1	7
M_3	2	2	4	5



$$MC_1(S) = 1 + 2 = 3$$



$$MC_2(S) = 1$$



$$MC_3(S) = 5$$

0

t

Is it a 1-envy free ? $\sum_{i \in S_1} p_{i1} = 1 + 2 = 3 \leq \sum_{i \in S_2} p_{i1} = 3$ M_1 does not envy M_2
 $\sum_{i \in S_1} p_{i1} = 1 + 2 = 3 \leq \sum_{i \in S_3} p_{i1} = 4$ M_1 does not envy M_3

It is easy to see that M_2 does not envy both M_1 and M_3

$$\sum_{i \in S_3} p_{i3} = 5 > \sum_{i \in S_1} p_{i3} = 2 + 2 = 4 \quad M_3 \text{ envies } M_1$$

It is not 1-envy-free

An example

- INPUT: 3 Unrelated machines, 4 jobs.

$p_{ih} =$

M \ J	J ₁	J ₂	J ₃	J ₄
M ₁	1	2	3	4
M ₂	3	5	1	7
M ₃	2	2	4	5



$$MC_1(S) = 1 + 2 = 3$$



$$MC_2(S) = 1$$



$$MC_3(S) = 5$$



Is it a 2-envy free ? $\sum_{i \in S_1} p_{i1} = 1 + 2 = 3 \leq 2 * \sum_{i \in S_2} p_{i1} = 2 * 3 = 6$ M₁ does not envy M₂

$\sum_{i \in S_1} p_{i1} = 1 + 2 = 3 \leq 2 * \sum_{i \in S_3} p_{i1} = 2 * 4 = 8$ M₁ does not envy M₃

It is easy to see that M₂ does not envy both M₁ and M₃

$\sum_{i \in S_3} p_{i3} = 5 \leq 2 * \sum_{i \in S_1} p_{i3} = 2 * 4 = 8$ M₃ does not envy M₁


Yes! it is 2-envy-free $\sum_{i \in S_3} p_{i3} = 5 \leq 2 * \sum_{i \in S_2} p_{i3} = 2 * 4 = 8$ M₃ does not envy M₂

An example

- INPUT: 3 Unrelated machines, 4 jobs.

$p_{ih} =$

$M \backslash J$	J_1	J_2	J_3	J_4
M_1	1	2	3	4
M_2	3	5	1	7
M_3	2	2	4	5




Since such schedule is 2-envy-free, then it is k -envy-free for any $k \geq 2$

An example

- **INPUT:** 3 Unrelated machines, 4 jobs.

$p_{ih} =$

$M \setminus J$	J_1	J_2	J_3	J_4
M_1	1	2	3	4
M_2	3	5	1	7
M_3	2	2	4	5



An 1-envy-free schedule!

Notice that a 1-envy-free solution (schedule) can be always obtained by assigning all jobs to one machine.

That is, k -envy-free solutions always exist! For any $k \geq 1$

Our goal

- We are interested in bounding the *Price of k -envy-freeness* that is the ratio between the MAKESPAN of the best k -envy-free schedule and the MAKESPAN of an optimal scheduling (non necessarily k -envy-free).
- The *Price of k -envy-freeness* is important because measures how much we lose by requiring that schedulings have to be k -Envy-free.

An upper bound to the Price of k-envy freeness ($k \geq 1$)

Theorem 11: The Price of k-envy-freeness for unrelated machines is at most $(1+1/k)^{\min\{n,m\}-1}$, for any $k \geq 1$.

Proof:

- We are going to show an algorithm that takes in input an optimal schedule OPT for the problem without k-envy-free constraint and transforms OPT into a k-envy-free scheduling whose MAKESPAN is at most $(1+1/k)^{\min\{n,m\}-1}$ the MAKESPAN of OPT .
- Let us consider the algorithm of the next slide:

...An upper bound to the Price of k-envy freeness ($k \geq 1$)

The algorithm:

INPUT: Optimal schedule OPT for the problem without k-envy-free constraint.

1. $S = OPT$;
2. While there exists a pair of machines (j, j') such that j envies j' then:
 - $S_j = S_j \cup S_{j'}$;
 - $S_{j'} = \emptyset$;
3. End while.
4. Return S .

- Since the algorithm removes all the envious, the returned schedule S is k-envy-free.

... An upper bound to the Price of k-envy freeness ($k \geq 1$)

- The algorithm iteratively moves the jobs from a machine j' to j whenever j envies j' ;
- It is obvious that, since the number of non-empty machines in S is at most $\min\{n, m\}$, the number of iterations is at most $\min\{n, m\} - 1$;
- Moreover notice that when j envies j' , we have that:

$$\sum_{i \in S_{j'}} p_{ij} < \frac{\sum_{i \in S_j} p_{ij}}{k}$$

- It follows that at each iteration, the current processing time of machine j , as a consequence of the movement of the job from j' to j , increase by a factor of at most $(1+1/k)$;
- It results that the MAKESPAN of the final schedule S obtained from the algorithm is at most $(1+1/k)^{\min\{n, m\}-1}$ times the MAKESPAN of the optimal solution OPT .

□

A lower bound to the Price of k-envy freeness ($k \geq 1$)

Theorem 12: The Price of k-envy-freeness for unrelated machines is at least $[1+1/(k+\varepsilon)]^{\min\{n,m\}-1}$, for any $k \geq 1$ and any (small) $\varepsilon > 0$.

Proof:

- It is possible to show an instance where $\frac{C_{\max}(k\text{-Envy_OPT})}{C_{\max}(OPT)} \geq (1 + \frac{1}{k+\varepsilon})^{\min\{n,m\}-1}$ for any small $\varepsilon > 0$.
- (no proof here! Of course I can provide the proof to interested students!)

Final considerations

- An important feature of the proofs we used to upper bound the *Price of k-envy-freeness* in the various cases (for both identical and unrelated machines) is that they rely on polynomial time algorithms constructing k-envy-free assignments of “reasonably” low MAKESPAN.
- In particular in the case of Theorem 8, we find a scheduling without the need of an initial scheduling in input.
- However in the cases of Theorem 9 and Theorem 11, given in input an optimal solution *OPT* that minimizes the MAKESPAN to the problem without the k-envy-free constraint, all the designed algorithms rearrange the allocations defined by *OPT* so as to obtain in polynomial time a k-envy-free assignment *S* for which we were able to show that:

$$\frac{C_{\max}(S)}{C_{\max}(OPT)} \leq r$$

where r is the upper bound that depends on the different cases we have considered. For instance, in Theorem 8, $r = \min \{n, m\}$. In Theorem 9, $r = 1 + 1/k$. Finally, in Theorem 11, $r = (1 + 1/k)^{\min\{n, m\} - 1}$.

Final considerations (2)

- Unfortunately we are not always able to compute OPT in polynomial time!
- However, all the designed algorithms we used in theorems 9 and 11 to upper bound the *Price of k-envy-freeness*, given in input any schedule N that could be not k-envy-free and with MAKESPAN equal to α times (for some $\alpha \geq 1$) the optimal MAKESPAN (that is $C_{\max}(N) = \alpha * C_{\max}(OPT)$), rearrange the allocation defined by N so as to obtain in polynomial time a k-envy-free assignment S such that:

$$\frac{C_{\max}(S)}{C_{\max}(N)} \leq r$$

where r is the upper bound we showed on the different cases we have considered.

- It follows that:
$$\frac{C_{\max}(S)}{C_{\max}(N)} = \frac{C_{\max}(S)}{\alpha * C_{\max}(OPT)} \leq r$$

- It implies that
$$\frac{C_{\max}(S)}{C_{\max}(OPT)} \leq \alpha * r$$

Final considerations (3)

- It means that for the cases where the problem is NP-Hard (recall that the problem of computing the minimum MAKESPAN is NP-Hard for identical machines and therefore for unrelated machines) we have polynomial time algorithms that compute k-envy-free solution with “reasonably” low MAKESPAN.
- Indeed we can use a polynomial time algorithm to compute an approximated scheduling without the k-envy-free constraint, and then we apply our upper bound to the price of k-envy-freeness algorithms (theorems 9 and 11) to return the final k-envy-free solution.
- Recall we have already seen (Theorem 5) the Algorithm *MAKESPAN* that is a $(4/3 - 1/3m)$ -approximation algorithm for the Scheduling Identical Machine MAKESPAN (minimization) problem.
- In the literature (we have not seen them in this course) there exists a PTAS (polynomial time approximation scheme) for the Scheduling Identical Machine MAKESPAN (minimization) problem, and a 2-approximation algorithm for the Scheduling unrelated Machine MAKESPAN (minimization) problem.