

DEEP LEARNING MIT 6.S191

MAR 20 2021

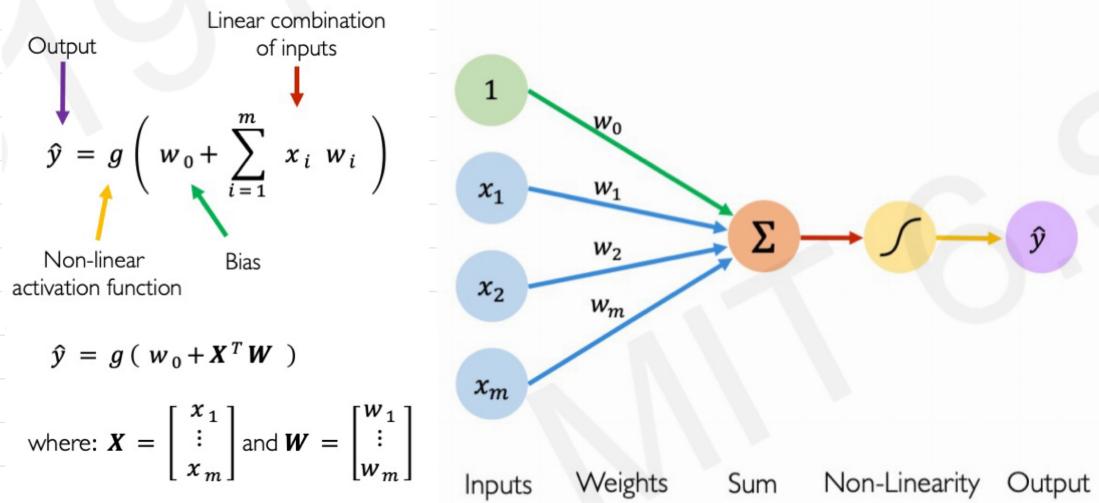


L1-INTRO To DL

BUILDING BLOCKS

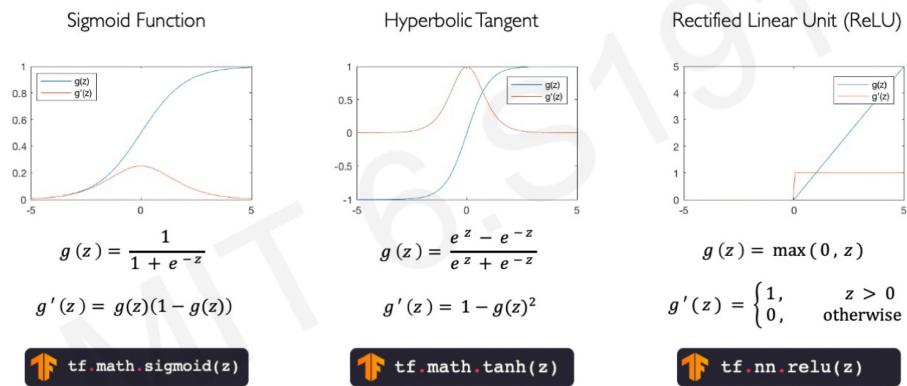
PERCEPTRON

FORWARD PROPAG.



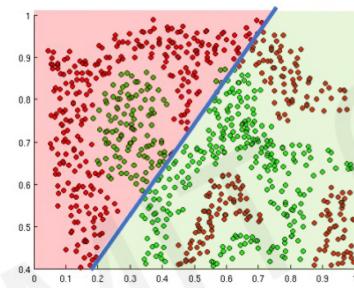
ACTIVATION FUNCTION

- Funzione (non lineare) che mappa gli output in un dato range

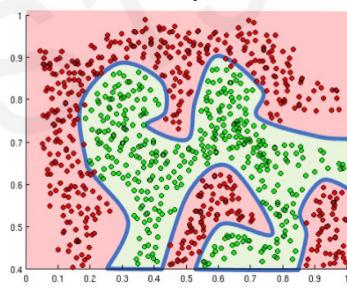


• La NON-LINEARITY ha lo scopo pratico di approssimare funzioni complesse

↳ Ad esempio per separare i punti colorati

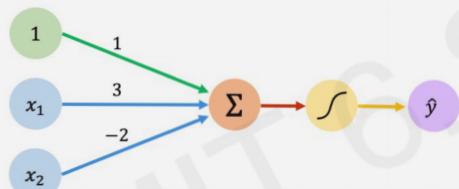


Linear activation functions produce linear decisions no matter the network size

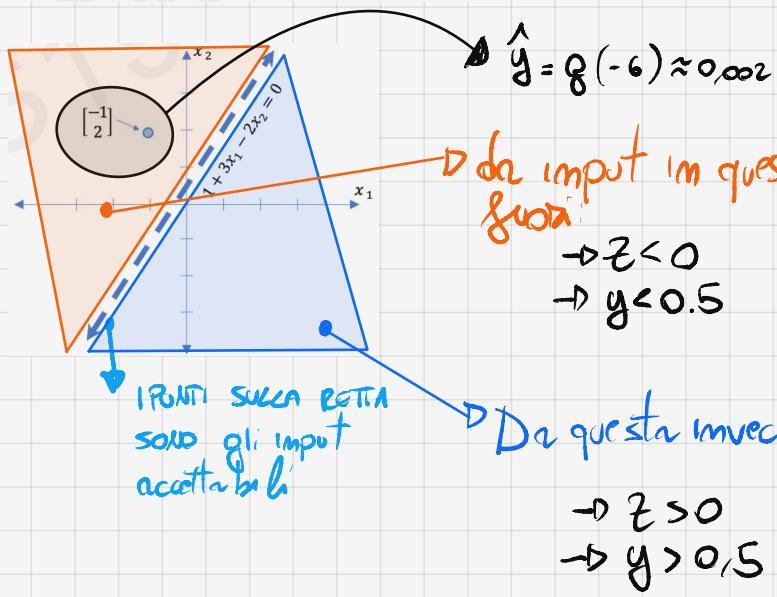


Non-linearities allow us to approximate arbitrarily complex functions

ESEMPIO



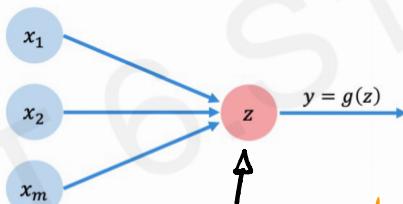
$$\hat{y} = g(1 + 3x_1 - 2x_2) \quad \text{2D}$$



From perceptron to Neural Net

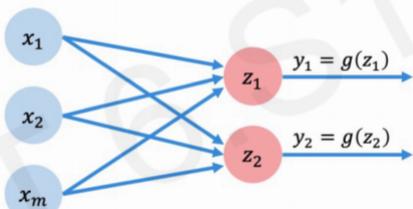
DENSE LAYER

• PERCEP. SEMPLIFICATE



Defn forward prop.

• MULTI-OUTPUT



→ Detto DENSE LAYER perché ogni z_i è comune a ogni x_j

In tensorflow

CLASS

```
class MyDenseLayer(tf.keras.layers.Layer):
    def __init__(self, input_dim, output_dim):
        super(MyDenseLayer, self).__init__()

        # Initialize weights and bias
        self.W = self.add_weight([input_dim, output_dim])
        self.b = self.add_weight([1, output_dim])

    def call(self, inputs):
        # Forward propagate the inputs
        z = tf.matmul(inputs, self.W) + self.b

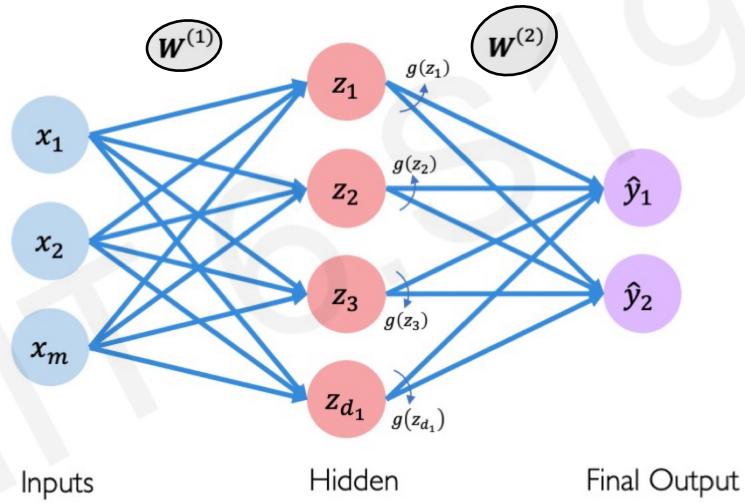
        # Feed through a non-linear activation
        output = tf.math.sigmoid(z)

        return output
```

INSTANCE

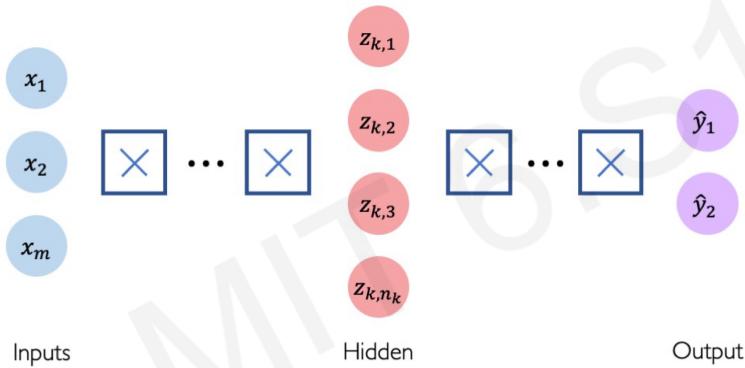
```
TF import tensorflow as tf
layer = tf.keras.layers.Dense(
    units=2)
```

SINGLE LAYER NN

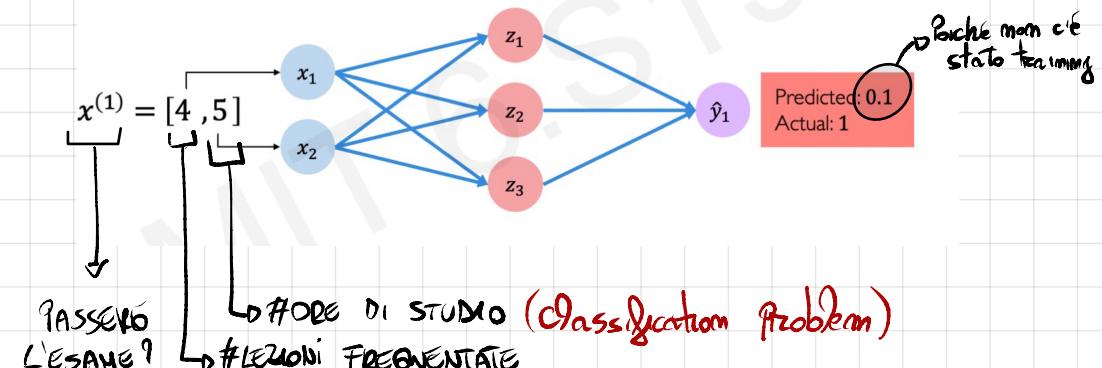


- Ogni layer ha una propria WEIGHT MATRIX

DEEP NN (multi-layer NN)



ESEMPIO

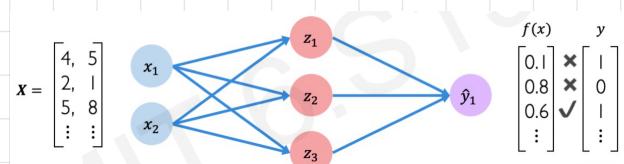


Come diciamo al modello quanto sbaglia?

LOSS

- $\mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$

Predicted Actual



Ld

Also known as:

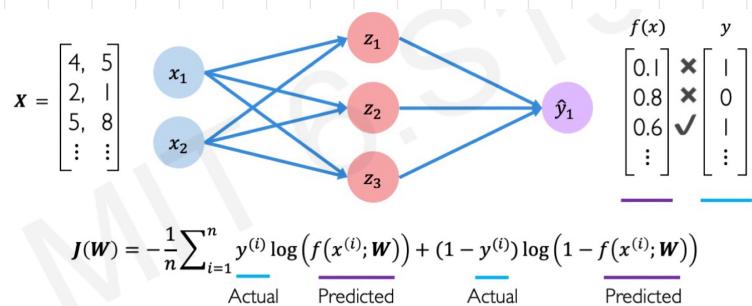
- Objective function
- Cost function
- Empirical Risk

$$J(\mathbf{W}) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$

Predicted Actual

EMPIRICAL LOSS

BINARY CROSS EMPTY LOSS

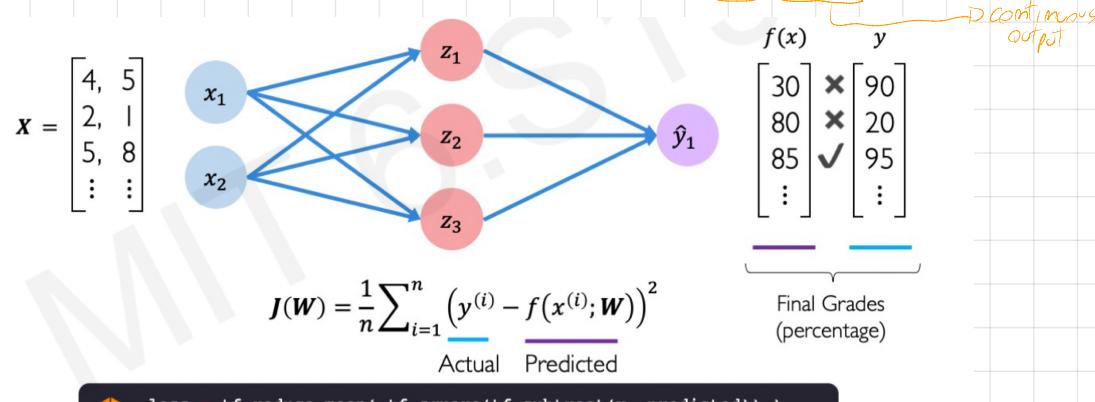


```
TensorFlow loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(y, predicted))
```

→ confronta la distribuzione di probabilità (ATTESA vs OTTENUTA)

MSE

- Se invece di un BINARY OUTPUT volessimo predire il VOTO, ci serve un'altra metrica per l'errore (regression problem)



```
TensorFlow loss = tf.reduce_mean(tf.square(tf.subtract(y, predicted)))  
loss = tf.keras.losses.MSE(y, predicted)
```

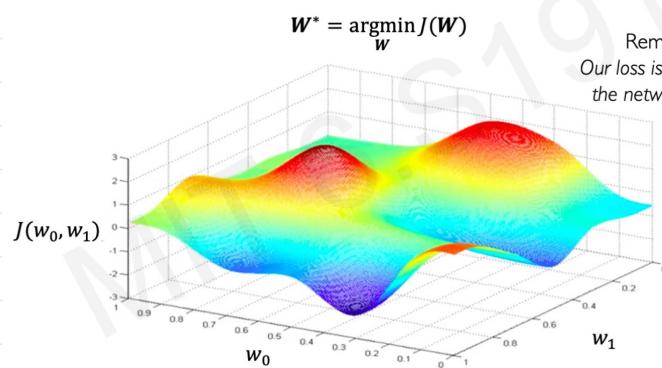
TRAINING

Loss Optimization

- Vogliamo i pesi che minimizzano la LOSS

$$\mathbf{W}^* = \underset{\mathbf{W}}{\operatorname{argmin}} J(\mathbf{W})$$

- Con un problema a due pesi, possiamo plottarne \mathbf{W}^*



es. GRADIENTE

- Scegliamo un punto nel grafico, calcoliamo i GRAD. e ci spostiamo nella sua direzione fino al LOCAL OPT.

GRADIENT DESCENT (ALGO)

NOTA se il learning rate è troppo piccolo restiamo bloccati in LOCAL OPT. se troppo grande si diffida la convergenza

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(W)}{\partial W}$
4. Update weights, $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

LEARNING RATE
(Quanto vogliamo imporre
in diretta)

```
import tensorflow as tf

weights = tf.Variable([tf.random.normal()])

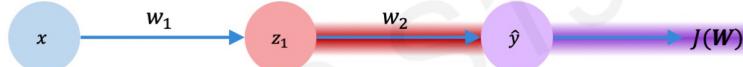
while True: # loop forever
    with tf.GradientTape() as g:
        loss = compute_loss(weights)
    gradient = g.gradient(loss, weights)

    weights = weights - lr * gradient
```

• Come calcoliamo il gradiente?

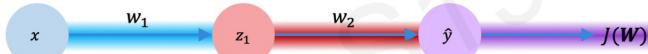
BACKPROPAGATION

- Ci interessa quanto un **PICCOLO CAMBIAMENTO** in un pes w_i influenzi $J(W)$
- Può essere espresso come rapporto tra derivate $\frac{\partial J(W)}{\partial w_2}$ e quale applichiamo la **CHAIN RULE**



$$\frac{\partial J(W)}{\partial w_2} = \frac{\partial J(W)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w_2}$$

• Possiamo espanderci a $\frac{\partial J(W)}{\partial w_1}$



$$\frac{\partial J(W)}{\partial w_1} = \frac{\partial J(W)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1} * \frac{\partial z_1}{\partial w_1}$$

- La chain rule si applica a tutti i modi ricorsivamente in NN con più layer

IN Practice

Come sceglio il LEARNING RATE?

ADAPTIVE LEARN. RATE

- Il GRADIENT DESCENT non funge con molti local opt.

- Il LEARNING RATE va settato per bene

↗ GRAD. DESC. ALGO

- Il LEARN. RATE non è fisso, ma è ottimizzato durante

- IL LEARNING PROCESS

Algorithm	TF Implementation
• SGD	tf.keras.optimizers.SGD
• Adam	tf.keras.optimizers.Adam
• Adadelta	tf.keras.optimizers.Adadelta
• Adagrad	tf.keras.optimizers.Adagrad
• RMSProp	tf.keras.optimizers.RMSProp

```

import tensorflow as tf
model = tf.keras.Sequential([...])
# pick your favorite optimizer
optimizer = tf.keras.optimizer.SGD()
while True: # loop forever
    # forward pass through the network
    prediction = model(x)
    with tf.GradientTape() as tape:
        # compute the loss
        loss = compute_loss(y, prediction)
    # update the weights using the gradient
    grads = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(grads, model.trainable_variables))

```



Can replace with any TensorFlow optimizer

LA BACKPROP è
possibile?

IDEA

- Impossibile com la BACKPROPAGATION è troppo da calcolare, soprattutto se lo facciamo per tutto il dataset.
- Lo computiamo per un sottoinsieme di punti

STOCHASTIC GRADIENT DESCENT (SGD)

- Ci permette di aumentare la learning rate

Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick single data point i
4. Compute gradient $\frac{\partial J_i(\mathbf{W})}{\partial \mathbf{W}}$
5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights

com un singolo punto è troppo noisy

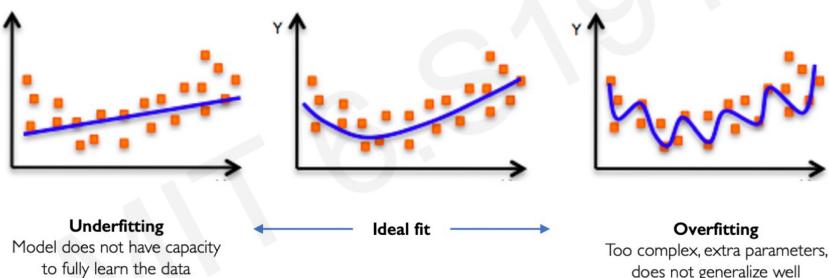
Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick batch of B data points
4. Compute gradient $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}} = \frac{1}{B} \sum_{k=1}^B \frac{\partial J_k(\mathbf{W})}{\partial \mathbf{W}}$
5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights

così va meglio

FITTING MODEL

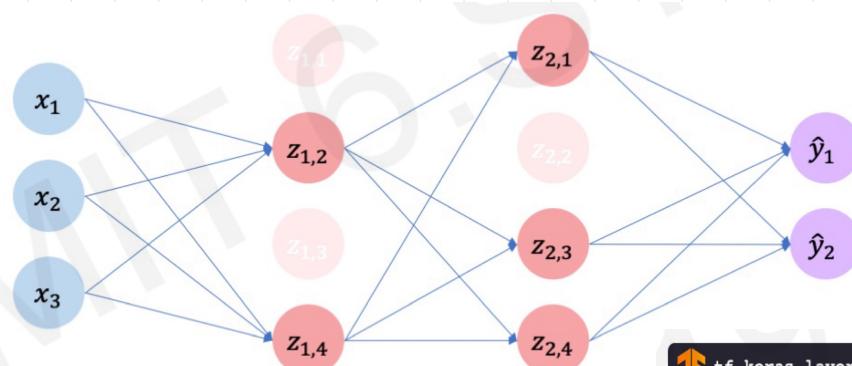
- Come rappresenta i dati il modello



- La linea blu è il modello usato per rappresentare i dati

REGULARIZATION DROPOUT

- Tecnica che impedisce al modello di andare in **OVERTFITTING**
- Durante il training settiamo alcune activation a \emptyset .

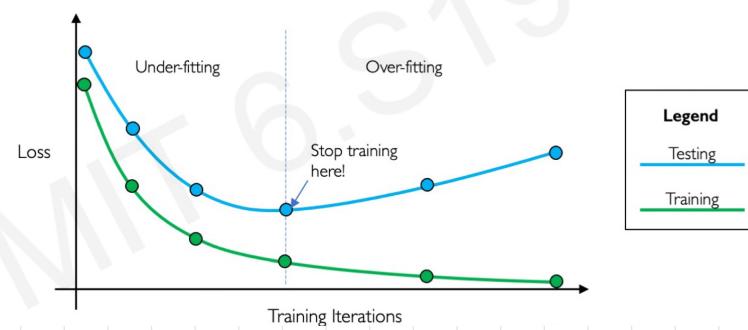


Lo indotte aumenta la velocità di train, poiché deve lavorare su meno modi.

EARLY STOPPING

- Smette il training quando risulta troppo loss

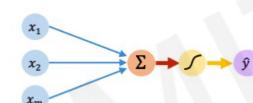
Stop training before we have a chance to overfit



COSA ABBIAMO FATTO?

The Perceptron

- Structural building blocks
- Nonlinear activation functions



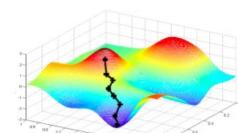
Neural Networks

- Stacking Perceptrons to form neural networks
- Optimization through backpropagation



Training in Practice

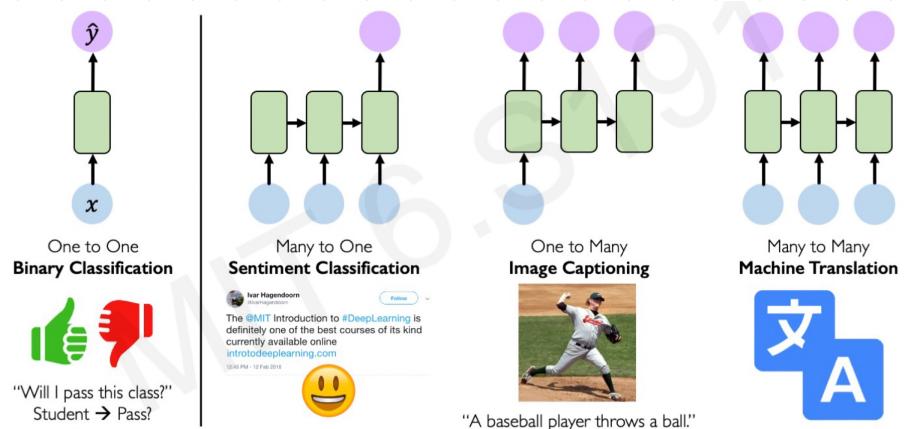
- Adaptive learning
- Batching
- Regularization



DEEP SEQUENCE MODELLING

- Lavoreremo su dati sequenziali nel tempo

Sequence modeling
Applications



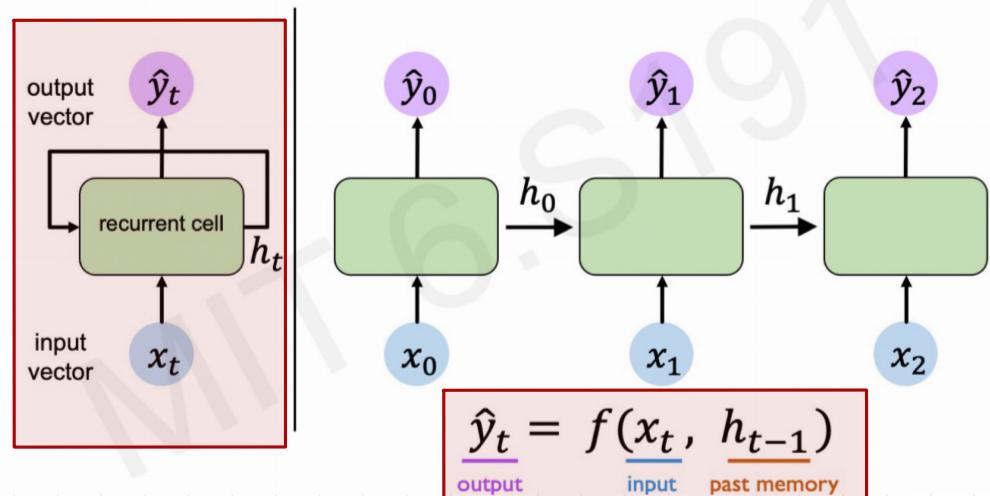
NEURONS WITH RECURRENCE

IDEA:

- i dati sequenziali sono **TIME STEPS** individuali, ci serve un modello che sia in grado di gestirli individualmente, ma anche prendendo in considerazione anche gli stati passati

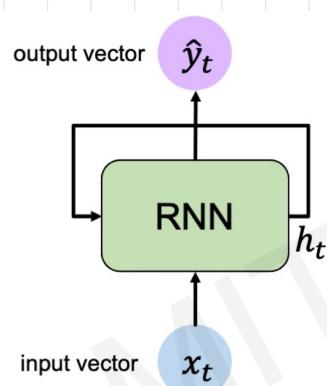
NEURONI CON
RICORRENZA

h_t sta per HIDDEN STATE



RECURRENT NEURAL NETWORKS (RNN)

RECURRENCE
RELATION



- Ad ogni time step calcoliamo il hidden state h_t associato

$$h_t = f_w(x_t, h_{t-1})$$

cell state function with weights w
input old state

```

my_rnn = RNN()
hidden_state = [0, 0, 0, 0]

sentence = ["I", "love", "recurrent", "neural"]

for word in sentence:
    prediction, hidden_state = my_rnn(word, hidden_state)

next_word_prediction = prediction
# >>> "networks!"

```

Passi del processo

NOTA: a somo 3
WEIGHT MATRICES

(1) INPUT VECTOR x_t

(2) UPDATE di h_t

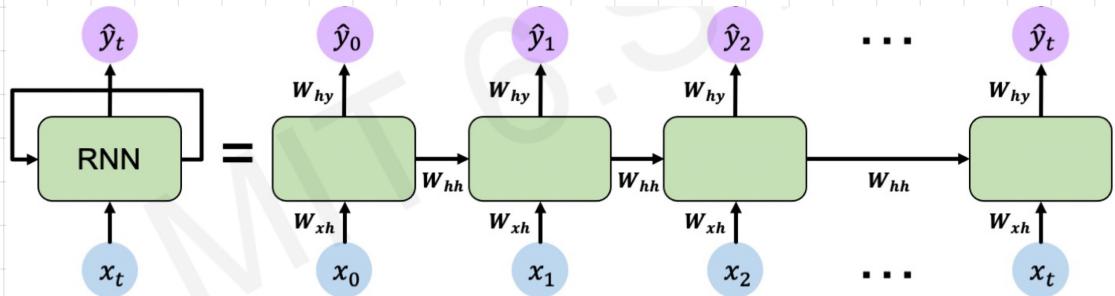
- moltiplichiamo h_{t-1} e x_t per le rispettive WEIGHT MATRICES

$$h_t = \tanh(W_{hh}^T h_{t-1} + W_{xh}^T x_t)$$

(3) CALCOLIAMO \hat{y}_t moltiplicando
il hidden state per un'altra
WEIGHT MATRIX

$$\hat{y}_t = W_{hy}^T h_t$$

... graficamente



OBSs

- Ad ogni output è associata una Loss, i cui insieme ci permette di allenare il modello.
- È importante da notare che le weight matrices sono sempre le stesse a ogni step.

oooPython

```
class MyRNNCell(tf.keras.layers.Layer):
    def __init__(self, rnn_units, input_dim, output_dim):
        super(MyRNNCell, self).__init__()

        # Initialize weight matrices
        self.W_xh = self.add_weight([rnn_units, input_dim])
        self.W_hh = self.add_weight([rnn_units, rnn_units])
        self.W_hy = self.add_weight([output_dim, rnn_units])

        # Initialize hidden state to zeros
        self.h = tf.zeros([rnn_units, 1])

    def call(self, x):
        # Update the hidden state
        self.h = tf.math.tanh( self.W_hh * self.h + self.W_xh * x )

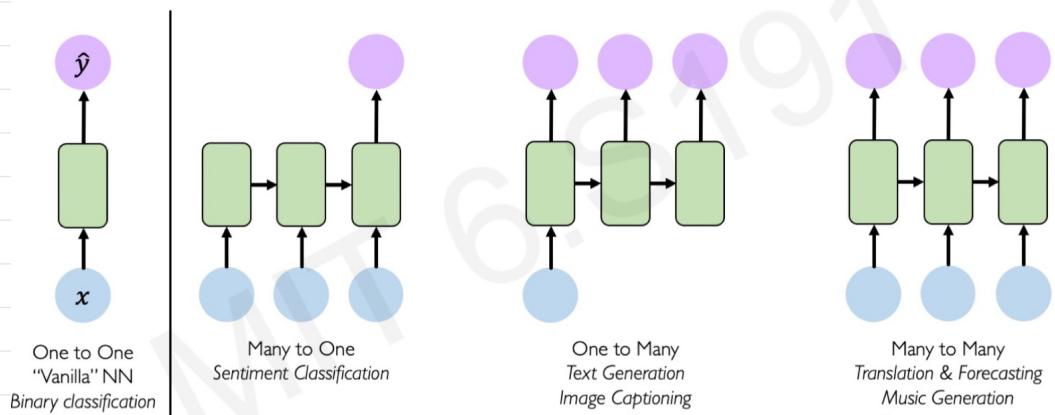
        # Compute the output
        output = self.W_hy * self.h

        # Return the current output and hidden state
        return output, self.h
```

- Ovviamente ci pensa tensorflow, non dobbiamo nulla implementarlo

```
tf.keras.layers.SimpleRNN(rnn_units)
```

RNNs per Sequential modeling



ENCODING ISSUE

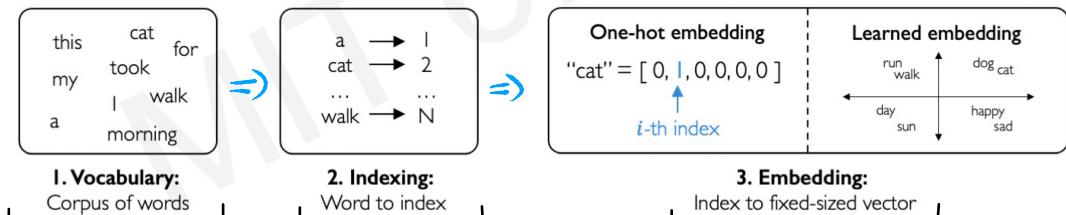
✗ "deep" ✗ "learning"

Neural networks cannot interpret words

✓ $\begin{bmatrix} 0.1 \\ 0.8 \\ 0.6 \end{bmatrix} \rightarrow \text{Learned embedding} \rightarrow \begin{bmatrix} 0.9 \\ 0.2 \\ 0.4 \end{bmatrix}$

Neural networks require numerical inputs

SOL



DESIGN CRITERIA

• 1) Modello deve avere l'abilità di:

(1) Gestire seq. di **LUNGHEZZA VARIABILE**

(2) Tracciare **LONG-TERM DEPENDENCIES**

"France is where I grew up, but I now live in Boston. I speak fluent ____."

(3) Mantenere informazioni sull'**ORDINE** → delle sequenze

(4) CONDIVIDERE PARAMETRI tra le sequenze

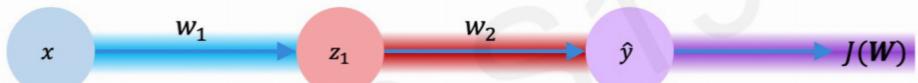
The food was great
vs.
We visited a restaurant for lunch
vs.
We were hungry but cleaned the house before eating

The food was good, not bad at all.
vs.
The food was bad, not good at all.
No food icon

BACKPROPAGATION THROUGH TIME (BPTT)

RIC

- Abbiamo visto che in BACKPROPAGATION consiste nel calcolare i GRADIENTI DELLA LOSS TOTAL ($\partial J(W)$) rispetto ogni parametro ($\frac{\partial J(W)}{\partial w_1}$), e poi modificare i parametri per diminuire la loss.

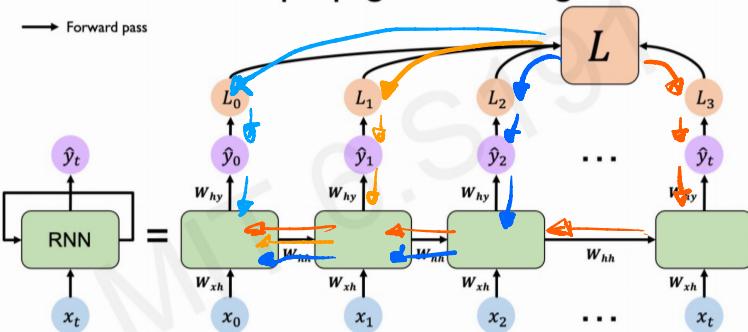


$$\frac{\partial J(W)}{\partial w_1} = \frac{\partial J(W)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1} * \frac{\partial z_1}{\partial w_1}$$

NON GRADIENT FLOW

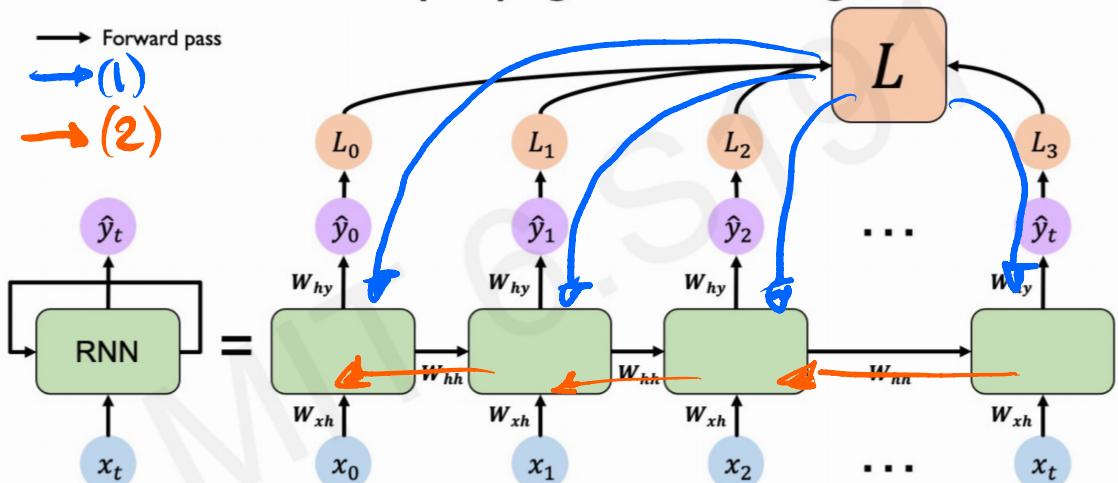
NO !

- Dala total loss invece di computare la Backprop. per ogni TIME STEP



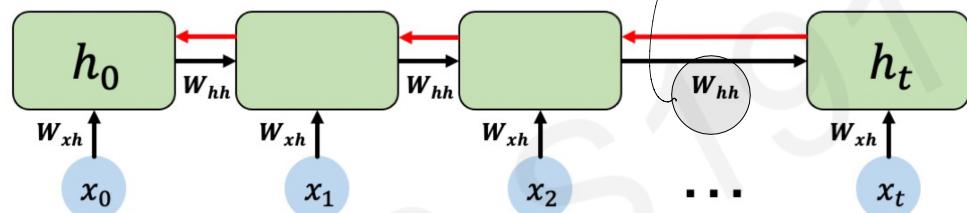
⇒ conviene farze un solo viaggio, dal TIME STEP più recente.
In questo modo l'errore ce lo potremo in una sola direzione.

SI !



- Calcolare il GRADIENTE rispetto h_0 , comporre:

- Molti moltiplicazioni per W_{hh}
- Computare molti gradienti



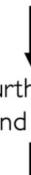
Exploding Gradient

Vanishing Gradient

- Se abbiamo molti valori > 1 in questo modo, i gradienti crescono ed overze moltiplicazioni per il GRADIENTE
- se abbiamo invece valori < 1 i GRADIENTI diminuiscono troppo

Why are vanishing gradients a problem?

Multiply many small numbers together

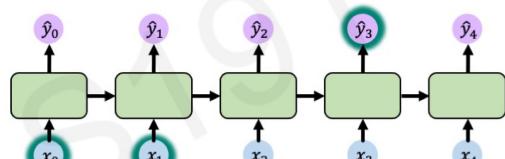


Errors due to further back time steps have smaller and smaller gradients

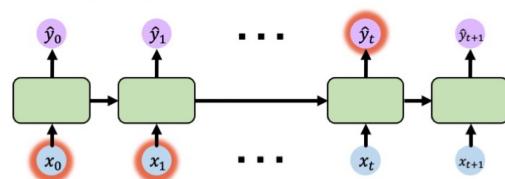


Bias parameters to capture short-term dependencies

"The clouds are in the ___"



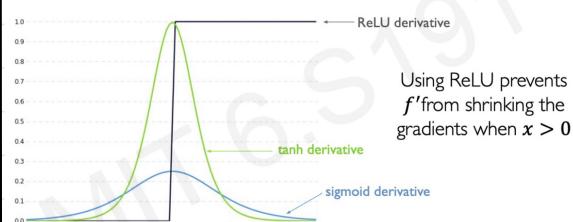
"I grew up in France, ... and I speak fluent ___"



SOL

- Ce sono vari trick per gestire i VANISHING GRADIENTS

Trick #1: Activation Functions



Using ReLU prevents f' from shrinking the gradients when $x > 0$

Trick #2: Parameter Initialization

Initialize **weights** to identity matrix

$$I_n = \begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{pmatrix}$$

Initialize **biases** to zero

This helps prevent the weights from shrinking to zero.

Solution #3: Gated Cells

Idea: use a more complex recurrent unit with gates to control what information is passed through

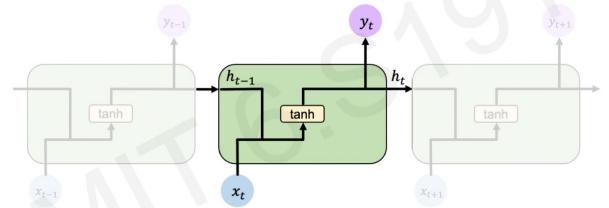
gated cell

LSTM, GRU, etc.

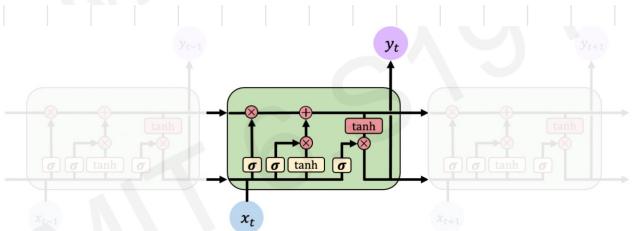
Long Short Term Memory (LSTMs) networks rely on a gated cell to track information throughout many time steps.

LONG SHORT TERM MEMORY

- STANDARD RNN



- LSTM RNN



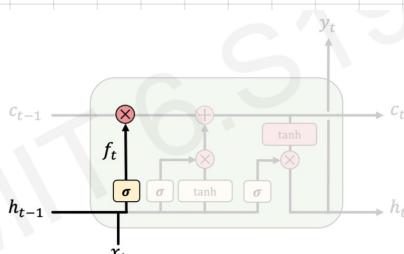
`tf.keras.layers.LSTM(num_units)`

- I GATES decidono quale informazione far passare o meno
- Lo fanno tramite, ad esempio, **SIGMOIDI** (come nelle imgs)

CELL STATE
(TIPPI) \rightarrow

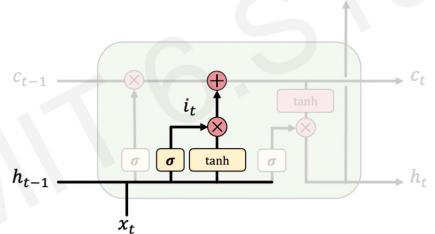
- Il parametro **C** è una variabile contenente informazioni passate
- Si divide in 4 fasi

(1)

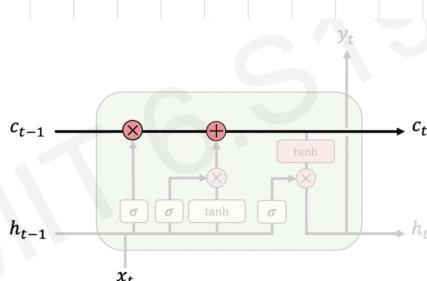


\rightarrow **FORGET**: butta parti irrilevanti dello stato precedente

(2) **STORE**: salva le info 1 irrellevanti sullo stato prossimo

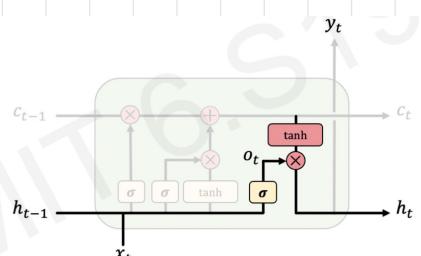


(3)

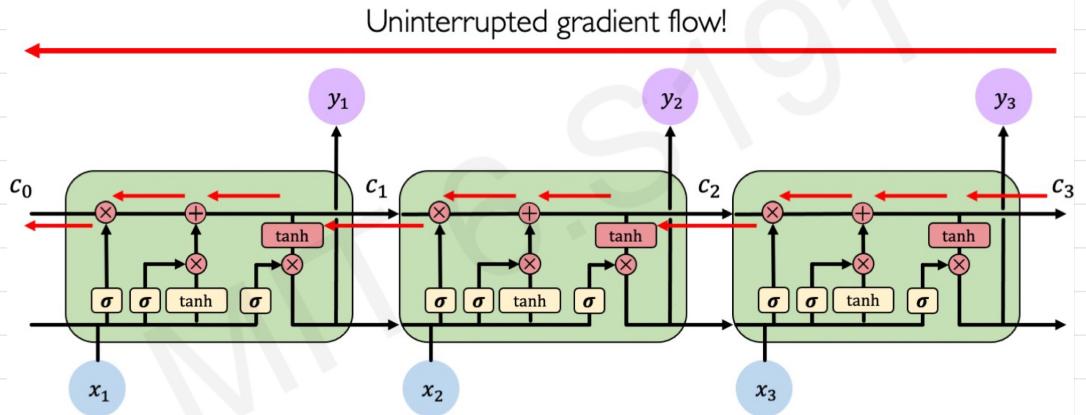


\rightarrow **UPDATE**: aggiorna il valore di C_t

(4) **OUTPUT**: controlla q'info \rightarrow invia alla stato successivo



• i9 GRADIENTE E ora calcolato attraverso i9 **CELL GATES** e
elimina i9 **VANISHING GRADIENT PROBLEM**



Concetti di LSTM

1. Maintain a **separate cell state** from what is outputted
2. Use **gates** to control the **flow of information**
 - **Forget** gate gets rid of irrelevant information
 - **Store** relevant information from current input
 - Selectively **update** cell state
 - **Output** gate returns a filtered version of the cell state
3. Backpropagation through time with **uninterrupted gradient flow**

Cosa abbiamo fatto?

Deep Learning for Sequence Modeling: Summary

1. RNNs are well suited for **sequence modeling** tasks
2. Model sequences via a **recurrence relation**
3. Training RNNs with **backpropagation through time**
4. Gated cells like **LSTMs** let us model **long-term dependencies**
5. Models for **music generation**, classification, machine translation, and more

DEEP COMPUTER VISION



157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	6	124	131	111	120	204	166	16	56	180
194	68	157	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	148	191	163	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	17	228	43	95	234
190	216	115	149	236	187	85	150	75	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	195	238	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	101	206	138	243	236
195	206	123	207	177	121	123	200	175	18	96	218

What the computer sees

157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	163	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	17	228	43	95	234
190	216	116	149	236	187	85	150	75	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	95	50	2	109	249	215
187	195	238	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	101	206	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

An image is just a matrix of numbers [0,255].
i.e., 1080x1080x3 for an RGB image

- Quali sono le task? - REGRESSION , CLASSIFICATION

CLASSIFICATION

- Si basa sull'identificazione delle caratteristiche chiave



Nose,
Eyes,
Mouth



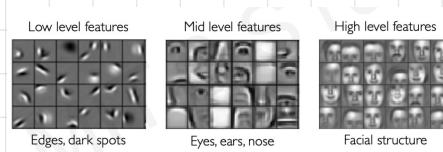
Wheels,
License Plate,
Headlights



Door,
Windows,
Steps

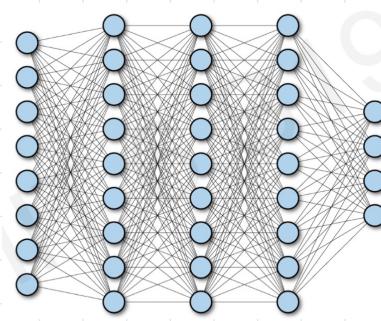
- È impossibile estrarre a mano le caratteristiche da classificare, ci sono troppe variazioni

- SIAMO IN GRADO DI ESTRARRE UNA GERARCHIA DI FEATURES DI RETTAMENTE DAI DATI?



LEARNING VISUAL FEATURE

- Abbriamo già visto come lavorare con FEED FORWARD, DENSE LAYER, RNN, ma qui c'è senso un'architettura diversa.



• Com'è un DENSE LAYER che prende in input un'IMMAGINE 2D, e fa perdere le informazioni spaziali (Invece di rimanere di gruppi di pixel ad esempio)

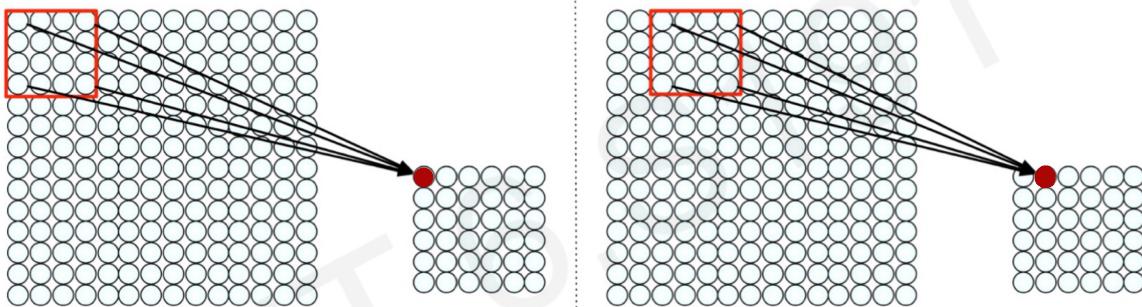


Questa operazione
è una **CONVOLUTION**



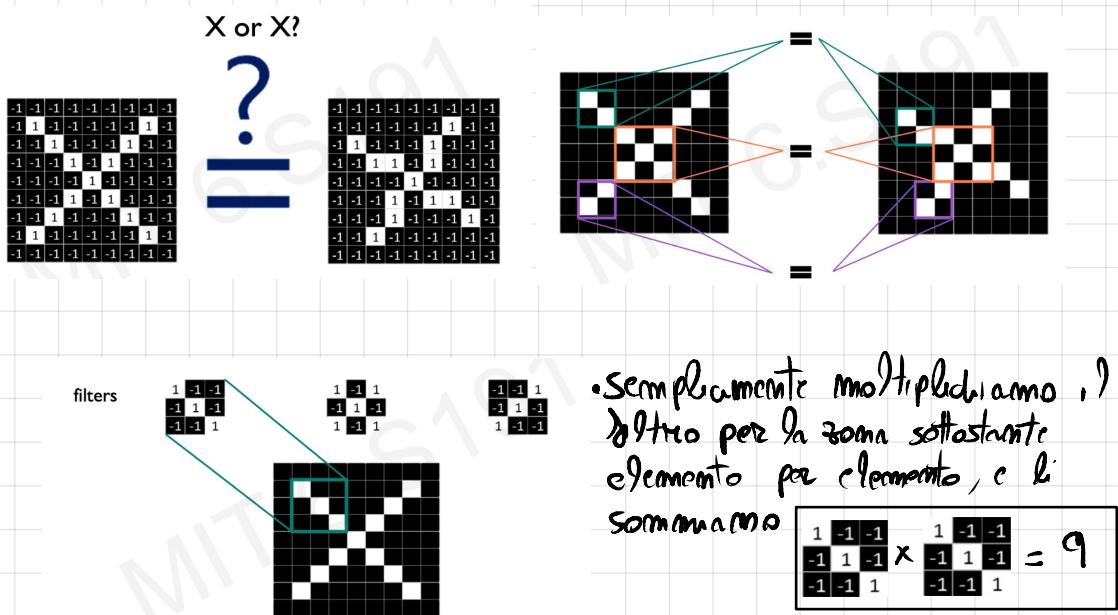
"Sfocare per far
risaltare caratteristiche
rilevanti"

- Utilizziamo un **SPATIAL STRUCTURE** e tramite un **MASCHERA** commettiamo un insieme di pixel nella stessa zona a un incisore in output, poi facciamo **SUDARE** la maschera



COME ASSEGNAANO I PESI ALLA MASCHERA?

ATTIVAZIONE DI
UN FILTRO



→ con lo slide del filtro, la matrice che esce dalle sue attivazioni è chiamata **FEATURE MAP**

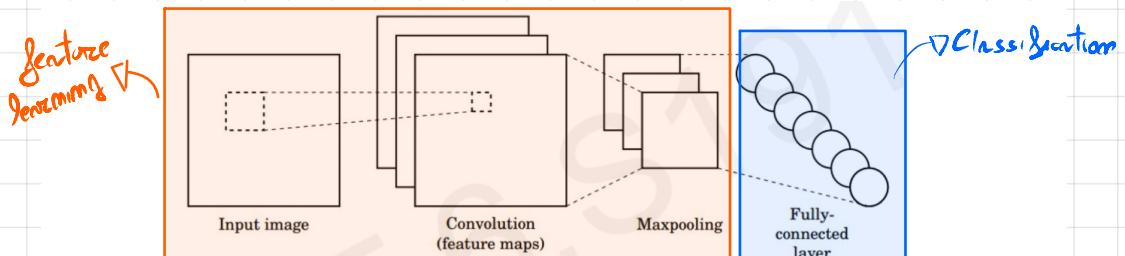


• A seconda del filtro
possiamo estrarre
le caratteristiche che
ci interessano

CONVENTIONAL NEURAL NETWORK (CNN)

CNN Per la CLASSIFICAZIONE

- Vogliamo estrarre le caratteristiche da un'immagine per poi mapparle su una **CLASSIFICATION TASK**
- Le operazioni di base sono:

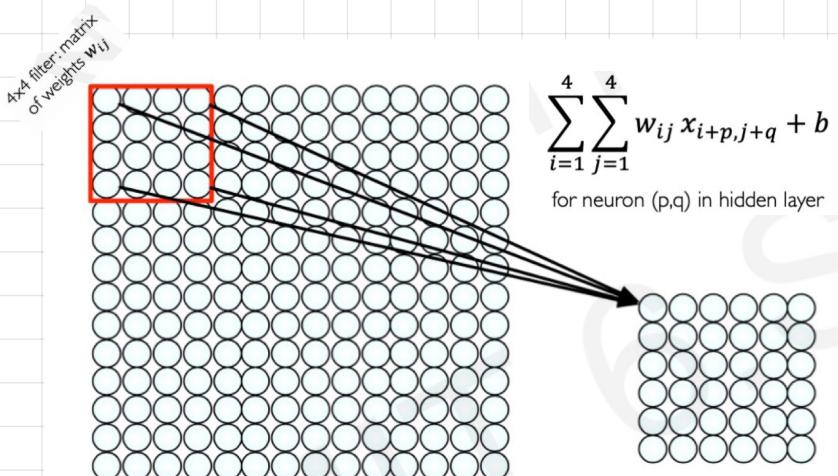


- Convolution:** Apply filters to generate **feature maps**.
- Non-linearity:** Often ReLU.
- Pooling:** Downsampling operation on each feature map.

`tf.keras.layers.Conv2D`
 `tf.keras.activations.*`
 `tf.keras.layers.MaxPool2D`

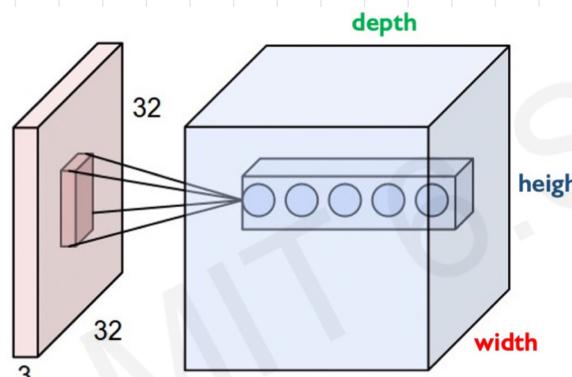
Lo si permette di ridurre la size delle feature map

CONVENTIONAL LAYER



COME POSSIAMO IN UN SINGOLO CONVENTIONAL LAYER
AVERE PIÙ FILTRI?

- Possiamo estrarre più caratteristiche da una stessa zona, applicando più filtri. Per ciò avremo un output con una terza dimensione: **DEPTH = #FILTRI**

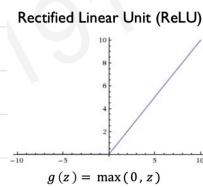


`tf.keras.layers.Conv2D(filters=d, kernel_size=(h,w), strides=s)`

ACTIVATION (RELU)

- Dobbiamo ricordarci l'activation function sulle feature map in output.

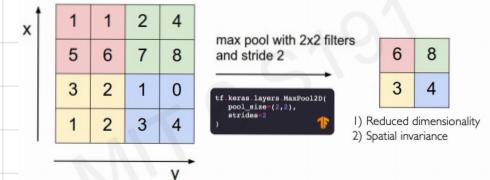
Inoltre si usano la funzione **NON LINEARE RELU** che porta tutti i valori negativi a 0.



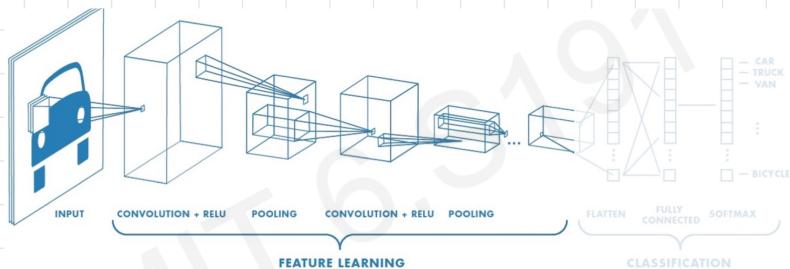
`tf.keras.layers.ReLU`

POOLING

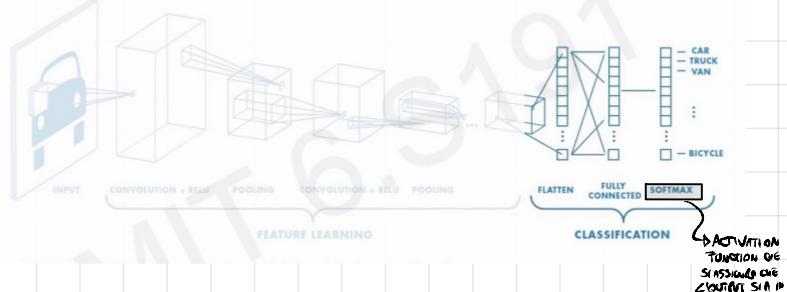
- Reduce la dimensione delle **FEAT MAP** preservando **SPATIAL INVARIANCE**



- Le operazioni descritte servono in sostanza per estrarre le feature **AD ALTO NIVELLO** e semplificare l'immagine al punto che può essere data in input a un **DENSE LAYER** senza i problemi descritti prima (**SPATIAL INVARIANCE**)



- Il **DENSE LAYER** a questo punto utilizza le feature per classificare l'immagine e il suo **OUTPUT** è in **PROB. DI APPARTENENZA A UNA DATA CLASSE**.



... in Tensorflow

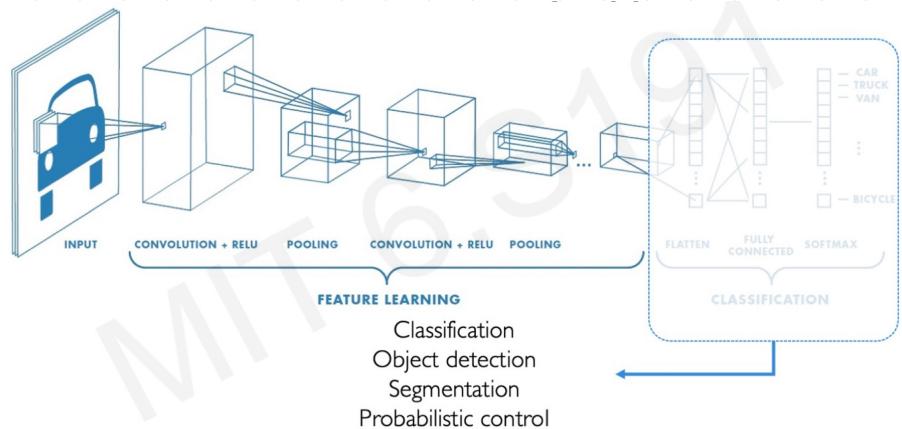
```
import tensorflow as tf

def generate_model():
    model = tf.keras.Sequential([
        # first convolutional layer
        tf.keras.layers.Conv2D(32, filter_size=3, activation='relu'),
        tf.keras.layers.MaxPool2D(pool_size=2, strides=2),
        # aumentato le caratteristiche da estrarre
        # second convolutional layer
        tf.keras.layers.Conv2D(64, filter_size=3, activation='relu'),
        tf.keras.layers.MaxPool2D(pool_size=2, strides=2),

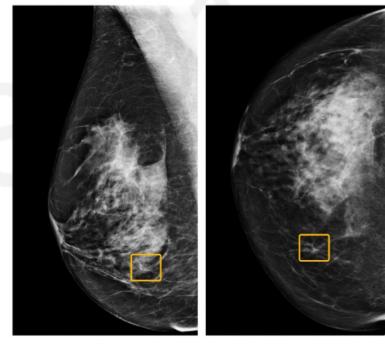
        # fully connected classifier
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(1024, activation='relu'),
        tf.keras.layers.Dense(10, activation='softmax') # 10 outputs
    ])
    return model
```

ALTRÉ APPLICAZIONI

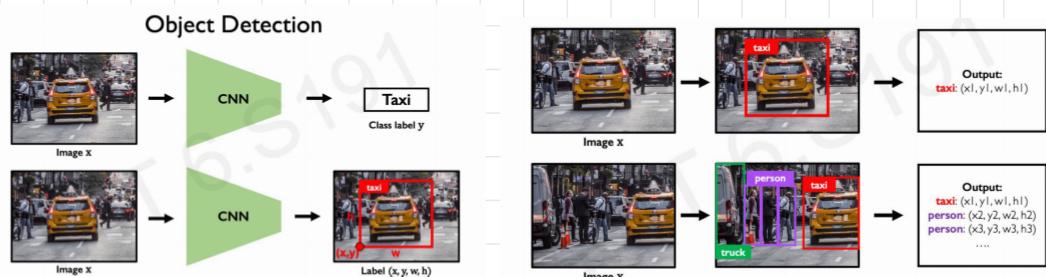
- I modelli definiti sopra non è limitato alla classificazione, ma a un'infinità di applicazioni



Esempio



OBJ DETECTION



Io a sì sono troppo dettagli!

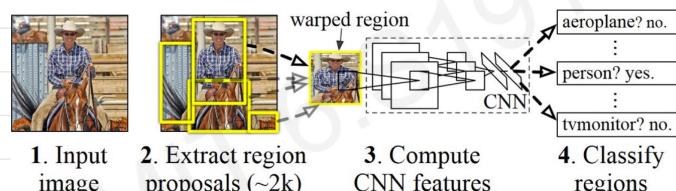
- Un approccio è di **SPARIRE BOX** sull'immagine e passare quelle sottoimmagini alla CNN.

Lo spartire spartire box a caso (o calcolarli tutti) è una **John**, ci serve qualche **HEURISTIC**

Object Detection with R-CNNs

①

R-CNN algorithm: Find regions that we think have objects. Use CNN to classify.



Problems: 1) Slow! Many regions; time intensive inference.

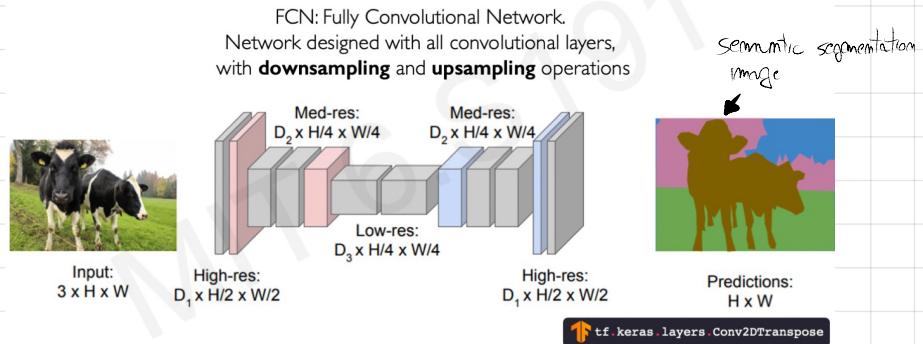
2) Brittle! Manually defined region proposals.

② Faster R-CNN Learns Region Proposals

VARIANTE del 'R-CNN': Al posto dell'algoritmo di prima inseriamo una parte del nostro modello TRAINED per riconoscere regioni interessanti della immagine

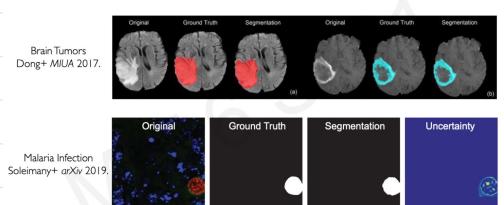
③

Semantic Segmentation: Fully Convolutional Networks

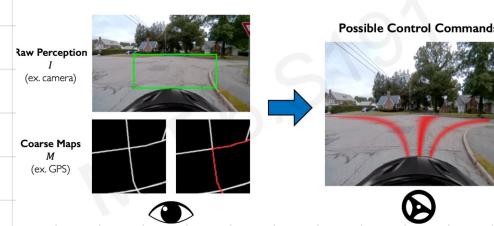


Aspettivamente classifica l'appartenenza di ogni pixel a una data regione, ad esempio i pixel delle mucche hanno tutta la stessa classe (colore)

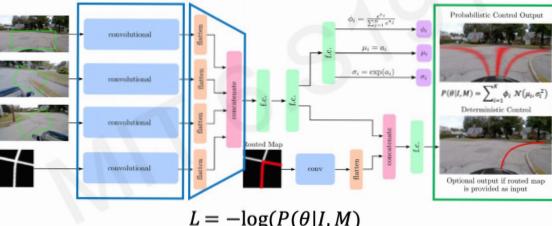
Semantic Segmentation: Biomedical Image Analysis



Continuous Control: Navigation from Vision



Entire model is trained end-to-end **without any human labelling or annotations**

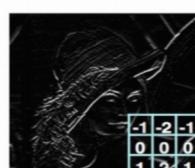


ooo SOMMARIO

Deep Learning for Computer Vision: Summary

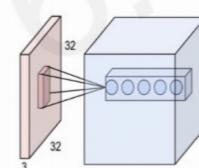
Foundations

- Why computer vision?
- Representing images
- Convolutions for feature extraction



CNNs

- CNN architecture
- Application to classification
- ImageNet



Applications

- Segmentation, image captioning, control
- Security, medicine, robotics

